

CS 432 – Interactive Computer Graphics

Lecture 4 – 3D Viewing

Reading

- Angel
 - Chapters 3-4
- Red Book
 - Chapter 5, Appendix E

From Objects to the Screen


- We started off by specifying objects via vertices that define primitives in 2D space.
 - These vertices were all in the range $-1 \leq x, y \leq 1$
- We added the ability to specify in OpenGL a *viewport* within the window to map this range to.
- In 3D space we can also specify a z-coordinate so that it's in view if $-1 \leq z \leq 1$

From Objects to the Screen

- In 2D space we call the area $-1 \leq x, y < 1$ the *clipping window*
 - Everything outside of this is clipped
- In 3D space, we call the *volume* $-1 \leq x, y, z < 1$ the *clipping volume*.
 - Everything outside of this volume is clipped

From Objects to the Screen

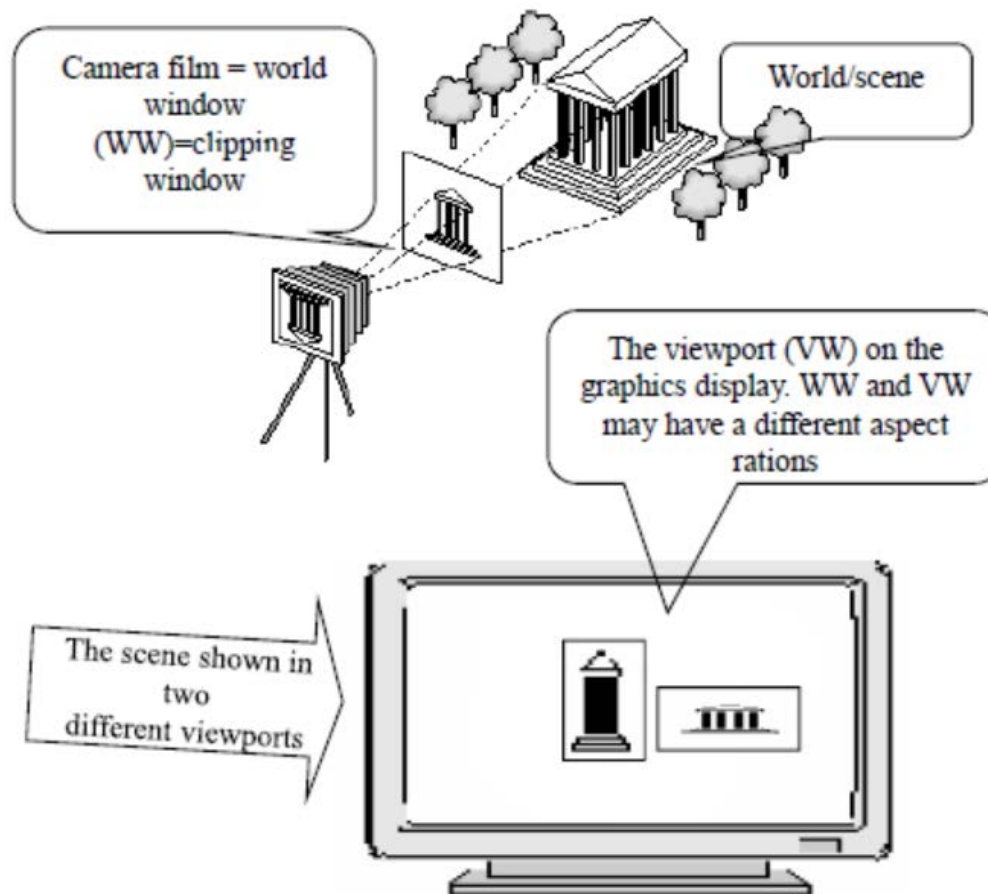
- Conceptually, the job of a virtual camera is to *project* points in the clipping volume onto the clipping area such that $z = 0$.
- These resulting points $(x, y, 0)$ are now considered in *camera coordinates*



From Objects to the Screen

- Once the objects are in camera coordinates they must be mapped to *screen coordinates* (i.e. pixels)
- We can even restrict what area of the screen to show
 - This is called the *viewport*

3D Viewing



OpenGL Coordinate Systems

- We now have four coordinate systems!

1. Model Coordinate System

- Where we build our objects

2. World Coordinate System


- Where we place our objects via a model matrix transformation

3. Camera Coordinate System

- Where the objects appear on the camera “film”


4. Window Coordinate System

- Where the objects on the film appear on the screen (pixels)



OpenGL Coordinate Systems

- Remember, when we were talking about forwards and backwards facing polygons we talked about the *right hand rule*.
- This also defines the orientation of these coordinate systems
- This result is that the positive z-axis points towards the viewer center-of-projection (COP)




Computer Viewing

- We can simulate viewing the world by transforming objects using the following pipeline:
 1. Position the objects
 - Via their respective model matrices
 2. Positioning the camera
 - Setting the view matrix
 3. Selecting the projection type/parameters
 - Setting the projection matrix
- Then we clip the transformed objects against some default camera view window or volume.

Default Coordinate Systems

- By default the object, world, and camera coordinate systems are the same resulting in the camera located at the shared origin and pointing in the negative z-direction
- The default camera view volume is a cube with sides of length 2 centered at the origin




Moving the Camera Frame

- If we want to visualize object with both positive and negative z values we can either
 - Move the camera in positive z direction
 - Translate the camera frame
 - Move the objects in the negative z direction
 - Translate the world frame
- Both of these view are equivalent and are determined by the model-view matrix
 - Want a translation: $Translate(0.0, 0.0, -d)$
 - Where $d > 0$


Moving Camera back from Origin

frames after translation by $-d$,
where $d > 0$

default frames




(a)



(b)

Moving the Camera

- We can move the camera to any desired position and orientation by a sequence of rotations and translations
- Example: side view
 - Rotate the camera
 - Move it away from the origin
 - Transformation matrix $M_{view} = TR$



OpenGL Code

- Remember that the last transformation specified is the first applied
 - `mat4 t = Translate(0.0, 0.0, -d);`
 - `mat4 ry = RotateY(90.0);`
 - `mat4 mView = t*ry;`
- So now we have both a model matrix for an object and a view matrix.
- We have options:
 - Send both independently to the shader and in the shader compute:
`gl_Position = view_matrix*model_matrix*vPosition;`
 - Compute `view_matrix*model_matrix` in the OpenGL program and send that single matrix to the shader program:
`modelview_matrix = view_matrix*model_matrix; //OpenGL`
`gl_Position = modelview_matrix*vPosition; //GLSL`
 - And of course you might just update the vertices right in OpenGL and send them to the shader.

Moving the Camera on the GPU


- Your choice may depend on how often things change
 - Never?
 - Model often, camera rarely?
 - Camera rarely, model often?
 - Both often?
- Remember, we most likely need to transpose our matrices when we send them to the shaders:

Transpose

```
GLuint model_view_loc =  
    glGetUniformLocation(program, "modelview_matrix");  
glUniformMatrix4fv(modelview_loc, 1, GL_TRUE, modelview_matrix);
```


Viewing APIs

- Specifying/computing the translation and rotation(s) may not be easy or natural.
- To simplify things, there are several APIs that will provide the model-view matrix for different parameters.
- One of the more popular (and supported in OpenGL) is the Look-At API
- But first we need to review some geometry!




Points

- We already know that to defined objects we need to provide vertex locations.
- Geometrically we can call these locations in space *points*.
- In graphics there are at least three other important geometric objects:
 1. Vectors
 2. Lines
 3. Planes

Vectors

- A vector is a quantity with two attributes
 - Direction
 - Magnitude
- Note: They have not actual position in space
- Examples include
 - Force
 - Velocity
 - Directed line segments
 - Important for graphics!




Vectors

- A vector can be defined given two points via subtraction:

$$v = P - Q$$


- Here we say the vector goes in the direction *from* point Q *to* point P .



Vectors




- Also given a point and a vector we can arrive at a new point via addition:

$$P = v + Q$$



Vector Operations

- There are additional properties of vectors:
 - Every vector has an inverse
 - Same magnitude but points in the opposite direction
 - Every vector can be multiplied by a scalar
 - Same direction (assuming non-negative magnitude), different magnitude
 - The sum of any two vectors is a vector
 - Use the head to tail axiom



Vectors in Homogenous Coordinates

- We know in 3D we specify vertices (points) via a 4D “vector” (confusing?) with the 4th coordinate equal to one:

$$P = (P_x, P_y, P_z, 1)$$

- Via vector arithmetic, we then define a vector in 3D homogenous coordinates as a 4D vector with the 4th coordinate equal to zero:

$$v = (v_x, v_y, v_z, 0)$$

Vectors in Homogenous Coordinates

$$v = (v_x, v_y, v_z, 0)$$

- The magnitude of this vector is:



$$\|v\| = \sqrt{v_x^2 + v_y^2 + v_z^2}$$

- Usually we specify the direction by making the vector have a length of one (unit length, or a normalized):

$$\frac{v}{\|v\|}$$

Planes


- A plane can be define either by
 - A point and two vectors
 - Three points




Normals

- Every plane has a vector n *normal* (perpendicular, orthogonal) to it
- Given two non-collinear vectors on a plane, u and v , we can use the **cross product** to find a plane's normal:

$$n = u \times v$$



Cross Product



- The cross product is defined as


$$a \times b = \|a\| \|b\| \sin(\theta) n$$

- Where n is the normal.
- So by solving for the cross product we get the normal multiplied by a scalar.
- Given two vectors u, v we can compute the cross product as

$$u \times v = ((u_y v_z - u_z v_y), (u_z v_x - u_x v_z), (u_x v_y - u_y v_x), 0)$$

- Fortunately both OpenGL and GLSL also give us a *cross* function
- **But be careful...**
 - Always make sure your 4th component is 0 so that it's treated as a vector!

Viewing APIs



- Ok back to a viewing API.....
- We can define a camera's location and orientation by specifying:
 - The position of the camera, often called the **eye** or the *view reference point* (VRP)
 - A vector that specifies the normal of the camera's view plane. This is often called the view plane normal (VPN, n in the figure below)
 - It is in the viewing direction In OpenGL the VPN is in the direction opposite the one in which the camera is looking
 - The direction opposite the VPN is often called the **at** vector.
 - Another vector, called the *view-up vector*, is in a direction specifying which is the approximate “**up**” direction for the camera


The LookAt Function

- The GLU library contained a function *gluLookAt* to form the required model-view matrix through a simple interface
- Replaced by `LookAt ()` in `mat.h`
 - These too will have to be transposed when sent to the shader programs.

```
mat4 view_matrix = LookAt(vec4 eye, vec4 at, vec4 up);
```


Angle.h Lookat

```
mat4 LookAt(const vec4& eye, const vec4& at, vec4& up){  
    vec4 n = normalize(eye-at);  
    vec4 u = normalize(cross(up,n));  
    vec4 v = normalize(cross(n,u));  
    vec4 t = vec4(0.0, 0.0, 0.0, 1.0);  
    mat4 c = mat4(u,v,n,t);  
    return c*Translate(-eye);  
}
```



The LookAt Function

- If we want to update our camera (move it, rotate it, etc..) we'll want to keep track of some things...
 - The location of the camera : eye
 - Three vectors that define the view plane:
 1. The view plane normal (opposite the at direction), n
 2. The "right" direction , u
 3. For convenience, we'll keep track of the third vector orthogonal to the other two (approximately the up direction) , v



The LookAt Function

- From these vectors you should easily be able to compute the *at* and *up* vectors needed for the LookAt API
 - $at = eye - n;$
 - $up = v;$

Other Viewing APIs


- The `LookAt` (and `gluLookAt`) function is one of many possible APIs for positioning the camera
- Others include
 - View reference point, view plane normal, view up
 - Yaw, pitch, roll
 - Elevation, azimuth, twist
 - Direction angles

Interface: Flying Cameras

- Often we like to use the keyboard to navigate around the “world”
- This is akin to moving the camera (and re-orienting it)
- If we use the Look-At API on each relevant keystroke we must
 - Compute the new eye point
 - Compute the new at point
 - Compute the new up vector

Interface: Flying Cameras

- It may be conceptually easier to move the camera according to another common API: pitch/yaw/roll
 - Pitch: rotate around current u axis
 - Roll: rotate around current n axis
 - Yaw: rotate around current v axis
- And update our coordinate system



Interface: Flying Cameras

- From our Look-At API we already have v (up) and $n=e-a$.
 - So $a=e-n$
 - And we can get $u=\text{normalize}(\text{cross}(v,n));$
- Pitch (rotate around u axis) by α radians
 - $u' = u$
 - $v' = \cos(\alpha)v - \sin(\alpha)n$
 - $n' = \sin(\alpha)v + \cos(\alpha)n$
- We can do this similarly for roll and yaw
- Notes:
 - Use the *normalize* function when necessary to get unit length vectors
 - **Again, be careful that your 4th component is zero before doing this.**
 - Make sure to call the `glutPostRedisplay()` after changing view matrix so that it is redrawn immediately

Interface: Flying Cameras

- The last thing we may want to do is move along (or away from) the direction of $-n$
- We can do this by change the eye and at points by the same amount along the n vector
 - $e = e + \beta n$
 - $a = a + \beta n$
- Interface:
 - $X \rightarrow$ Pitch down
 - $x \rightarrow$ Pitch up
 - $C \rightarrow$ Yaw clockwise in the uv plane
 - $c \rightarrow$ Yaw counter-clockwise in the uv plane
 - $Z \rightarrow$ Roll clockwise in the uv plane
 - $z \rightarrow$ Roll counter-clockwise in the uv plane
 - `GLUT_KEY_UP` \rightarrow Move towards at point
 - `GLUT_KEY_Down` \rightarrow Move away from at point
- NOTE: `GLUT_KEY_UP/DOWN` are “special keys” and are called via the `glutSpecialFunc` callback

A Camera Class

- For convenience you'll likely want to create a Camera class.
- This can store:
 - Current `view_matrix`
 - `eye`
 - `u, v, n`
- And can do things like
 - Get the `view_matrix`
 - Update `eye, u, v, n` based on interactions.
- And we'll add more (the projection matrix) in a minute...
- And when it comes time to draw an object, just pass it the current `view_matrix` so it can either use it or pass it to the shader.